

CSP 571 Project

Cab Ride Data Analytics

A20489131 - Pranjal Naik

A20488668 - Kishan Patel

A20468078 - Jainam Shah

Loading Libraries

```
#install.packages("caret", repos = "http://cran.us.r-project.org")
# install.packages("sqldf")
# install.packages("tidyr")
# install.packages("tidyverse")
# install.packages("ggplot2")
# install.packages("readr")
#install.packages("gmodels")
# install.packages("tm")
# install.packages("SnowballC")
# install.packages("wordcloud")
# install.packages("RColorBrewer")
# install.packages("treemap")
# install.packages("highcharter")
# install.packages("remotes")
# remotes::install_github("cran/DMwR")
#install.packages("corrplot")
#install.packages('rpart.plot')
library(rpart.plot)
```

```
## Loading required package: rpart
```

```
library(corrplot)
```

```
## corrplot 0.92 loaded
```

```
library("DMwR")
```

```
## Loading required package: lattice
```

```
## Loading required package: grid
```

```
## Registered S3 method overwritten by 'quantmod':
##   method           from
##   as.zoo.data.frame zoo
```

```
library(treemap)
library(highcharter)
library(tidyr)
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
##
##   filter, lag
```

```
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library(tidyverse)
```

```
## — Attaching packages
## _____
## tidyverse 1.3.2 —
```

```
## ✓ ggplot2 3.4.0      ✓ purrr 0.3.5
## ✓ tibble 3.1.8       ✓ stringr 1.4.1
## ✓ readr 2.1.3        ✓ forcats 0.5.2
## — Conflicts ————— tidyverse_conflicts() —
## ✖ dplyr::filter() masks stats::filter()
## ✖ dplyr::lag() masks stats::lag()
```

```
options(gsubfn.engine="R")
library(sqldf)
```

```
## Loading required package: gsubfn
## Loading required package: proto
## Loading required package: RSQLite
```

```
library(ggplot2)
library(readr)
library(gmodels)
library(tm)
```

```
## Loading required package: NLP
##
## Attaching package: 'NLP'
##
## The following object is masked from 'package:ggplot2':
##
##      annotate
```

```
library(SnowballC)
library(wordcloud)
```

```
## Loading required package: RColorBrewer
```

```
library(RColorBrewer)
library(rpart)
library(randomForest)
```

```
## randomForest 4.7-1.1
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
##
## The following object is masked from 'package:ggplot2':
##
##      margin
##
## The following object is masked from 'package:dplyr':
##
##      combine
```

Loading Dataset

```
cab_ds <- read.csv("/Users/Shared/pranjalnaik/Pranjal DPA/Dataset/rideshare_kaggle.csv")
#head(cab_ds)
summary(cab_ds)
```

```

##      id      timestamp      hour      day
## Length:693071   Min.   :1.543e+09   Min.   : 0.00   Min.   : 1.00
## Class :character 1st Qu.:1.543e+09   1st Qu.: 6.00   1st Qu.:13.00
## Mode  :character Median :1.544e+09   Median :12.00   Median :17.00
##              Mean  :1.544e+09   Mean  :11.62   Mean  :17.79
##              3rd Qu.:1.545e+09   3rd Qu.:18.00   3rd Qu.:28.00
##              Max.   :1.545e+09   Max.   :23.00   Max.   :30.00
##
##      month      datetime      timezone      source
## Min.   :11.00   Length:693071   Length:693071   Length:693071
## 1st Qu.:11.00   Class :character  Class :character  Class :character
## Median :12.00   Mode  :character  Mode  :character  Mode  :character
## Mean    :11.59
## 3rd Qu.:12.00
## Max.    :12.00
##
## destination      cab_type      product_id      name
## Length:693071     Length:693071     Length:693071     Length:693071
## Class :character   Class :character   Class :character   Class :character
## Mode  :character   Mode  :character   Mode  :character   Mode  :character
##
##
##
##      price      distance      surge_multiplier      latitude
## Min.   : 2.50   Min.   :0.020   Min.   :1.000   Min.   :42.21
## 1st Qu.: 9.00   1st Qu.:1.280   1st Qu.:1.000   1st Qu.:42.35
## Median :13.50   Median :2.160   Median :1.000   Median :42.35
## Mean    :16.55   Mean    :2.189   Mean    :1.014   Mean    :42.34
## 3rd Qu.:22.50   3rd Qu.:2.920   3rd Qu.:1.000   3rd Qu.:42.36
## Max.    :97.50   Max.    :7.860   Max.    :3.000   Max.    :42.37
## NA's    :55095
##
## longitude      temperature      apparentTemperature      short_summary
## Min.   : -71.11   Min.   :18.91   Min.   :12.13   Length:693071
## 1st Qu.: -71.08   1st Qu.:36.45   1st Qu.:31.91   Class :character
## Median : -71.06   Median :40.49   Median :35.90   Mode  :character
## Mean    : -71.07   Mean    :39.58   Mean    :35.88
## 3rd Qu.: -71.05   3rd Qu.:43.58   3rd Qu.:40.08
## Max.    : -71.03   Max.    :57.22   Max.    :57.22
##
##
## long_summary      precipIntensity      precipProbability      humidity
## Length:693071     Min.   :0.000000   Min.   :0.0000   Min.   :0.3800
## Class :character  1st Qu.:0.000000   1st Qu.:0.0000   1st Qu.:0.6400
## Mode  :character  Median :0.000000   Median :0.0000   Median :0.7100
##              Mean  :0.008922   Mean  :0.1461   Mean  :0.7411
##              3rd Qu.:0.000000   3rd Qu.:0.0000   3rd Qu.:0.8800
##              Max.   :0.144700   Max.   :1.0000   Max.   :0.9600
##
##
## windSpeed      windGust      windGustTime      visibility
## Min.   : 0.450   Min.   : 0.80   Min.   :1.543e+09   Min.   : 0.717
## 1st Qu.: 3.410   1st Qu.: 4.06   1st Qu.:1.543e+09   1st Qu.: 8.432
## Median : 5.910   Median : 7.55   Median :1.544e+09   Median : 9.880

```

```

## Mean      : 6.186      Mean      : 8.47      Mean      :1.544e+09      Mean      : 8.468
## 3rd Qu.: 8.410      3rd Qu.:11.74      3rd Qu.:1.545e+09      3rd Qu.: 9.996
## Max.      :15.000      Max.      :27.25      Max.      :1.545e+09      Max.      :10.000
##
## temperatureHigh temperatureHighTime temperatureLow temperatureLowTime
## Min.      :32.68      Min.      :1.543e+09      Min.      :17.85      Min.      :1.543e+09
## 1st Qu.:42.57      1st Qu.:1.543e+09      1st Qu.:30.17      1st Qu.:1.543e+09
## Median :44.68      Median :1.544e+09      Median :34.18      Median :1.544e+09
## Mean      :45.04      Mean      :1.544e+09      Mean      :34.15      Mean      :1.544e+09
## 3rd Qu.:46.91      3rd Qu.:1.545e+09      3rd Qu.:38.73      3rd Qu.:1.545e+09
## Max.      :57.87      Max.      :1.545e+09      Max.      :46.60      Max.      :1.545e+09
##
## apparentTemperatureHigh apparentTemperatureHighTime apparentTemperatureLow
## Min.      :22.62      Min.      :1.543e+09      Min.      :11.81
## 1st Qu.:36.57      1st Qu.:1.543e+09      1st Qu.:27.70
## Median :40.95      Median :1.544e+09      Median :30.03
## Mean      :41.61      Mean      :1.544e+09      Mean      :30.14
## 3rd Qu.:44.12      3rd Qu.:1.545e+09      3rd Qu.:35.32
## Max.      :57.20      Max.      :1.545e+09      Max.      :47.25
##
## apparentTemperatureLowTime      icon      dewPoint      pressure
## Min.      :1.543e+09      Length:693071      Min.      : 4.39      Min.      : 988.1
## 1st Qu.:1.543e+09      Class :character      1st Qu.:27.49      1st Qu.: 999.8
## Median :1.544e+09      Mode  :character      Median :30.69      Median :1009.2
## Mean      :1.544e+09      Mean      :31.66      Mean      :1010.1
## 3rd Qu.:1.545e+09      3rd Qu.:38.12      3rd Qu.:1021.9
## Max.      :1.545e+09      Max.      :50.67      Max.      :1035.5
##
## windBearing      cloudCover      uvIndex      visibility.1
## Min.      : 2.0      Min.      :0.0000      Min.      :0.0000      Min.      : 0.717
## 1st Qu.:124.0      1st Qu.:0.3700      1st Qu.:0.0000      1st Qu.: 8.432
## Median :258.0      Median :0.8200      Median :0.0000      Median : 9.880
## Mean      :220.1      Mean      :0.6865      Mean      :0.2489      Mean      : 8.468
## 3rd Qu.:303.0      3rd Qu.:1.0000      3rd Qu.:0.0000      3rd Qu.: 9.996
## Max.      :356.0      Max.      :1.0000      Max.      :2.0000      Max.      :10.000
##
## ozone      sunriseTime      sunsetTime      moonPhase
## Min.      :269.4      Min.      :1.543e+09      Min.      :1.543e+09      Min.      :0.0900
## 1st Qu.:290.9      1st Qu.:1.543e+09      1st Qu.:1.543e+09      1st Qu.:0.3000
## Median :307.4      Median :1.544e+09      Median :1.544e+09      Median :0.6800
## Mean      :313.5      Mean      :1.544e+09      Mean      :1.544e+09      Mean      :0.5791
## 3rd Qu.:331.8      3rd Qu.:1.545e+09      3rd Qu.:1.545e+09      3rd Qu.:0.7900
## Max.      :378.9      Max.      :1.545e+09      Max.      :1.545e+09      Max.      :0.9300
##
## precipIntensityMax uvIndexTime      temperatureMin temperatureMinTime
## Min.      :0.00000      Min.      :1.543e+09      Min.      :15.63      Min.      :1.543e+09
## 1st Qu.:0.00000      1st Qu.:1.543e+09      1st Qu.:30.17      1st Qu.:1.543e+09
## Median :0.00040      Median :1.544e+09      Median :34.24      Median :1.544e+09
## Mean      :0.03737      Mean      :1.544e+09      Mean      :33.46      Mean      :1.544e+09
## 3rd Qu.:0.09160      3rd Qu.:1.545e+09      3rd Qu.:38.88      3rd Qu.:1.545e+09
## Max.      :0.14590      Max.      :1.545e+09      Max.      :43.10      Max.      :1.545e+09
##

```

```
## temperatureMax temperatureMaxTime apparentTemperatureMin
## Min. :33.51 Min. :1.543e+09 Min. :11.81
## 1st Qu.:42.57 1st Qu.:1.543e+09 1st Qu.:27.76
## Median :44.68 Median :1.544e+09 Median :30.13
## Mean :45.26 Mean :1.544e+09 Mean :29.73
## 3rd Qu.:46.91 3rd Qu.:1.545e+09 3rd Qu.:35.71
## Max. :57.87 Max. :1.545e+09 Max. :40.05
##
## apparentTemperatureMinTime apparentTemperatureMax apparentTemperatureMaxTime
## Min. :1.543e+09 Min. :28.95 Min. :1.543e+09
## 1st Qu.:1.543e+09 1st Qu.:36.57 1st Qu.:1.543e+09
## Median :1.544e+09 Median :40.95 Median :1.544e+09
## Mean :1.544e+09 Mean :42.00 Mean :1.544e+09
## 3rd Qu.:1.545e+09 3rd Qu.:44.12 3rd Qu.:1.545e+09
## Max. :1.545e+09 Max. :57.20 Max. :1.545e+09
##
```

Data Preprocessing

###Performing Data Sanity Checks before proceeding with analysis

```
#Checking the shape of the dataset
row=nrow(cab_ds)
col=ncol(cab_ds)
sprintf("The rows and columns are: %s %s",row,col)
```

```
## [1] "The rows and columns are: 693071 57"
```

```
#See whether missing values or not
sapply(cab_ds, function(x) sum(is.na(x)))
```

```

##          id          timestamp
##          0          0
##          hour          day
##          0          0
##          month          datetime
##          0          0
##          timezone          source
##          0          0
##          destination          cab_type
##          0          0
##          product_id          name
##          0          0
##          price          distance
##          55095          0
##          surge_multiplier          latitude
##          0          0
##          longitude          temperature
##          0          0
##          apparentTemperature          short_summary
##          0          0
##          long_summary          precipIntensity
##          0          0
##          precipProbability          humidity
##          0          0
##          windSpeed          windGust
##          0          0
##          windGustTime          visibility
##          0          0
##          temperatureHigh          temperatureHighTime
##          0          0
##          temperatureLow          temperatureLowTime
##          0          0
##          apparentTemperatureHigh apparentTemperatureHighTime
##          0          0
##          apparentTemperatureLow apparentTemperatureLowTime
##          0          0
##          icon          dewPoint
##          0          0
##          pressure          windBearing
##          0          0
##          cloudCover          uvIndex
##          0          0
##          visibility.1          ozone
##          0          0
##          sunriseTime          sunsetTime
##          0          0
##          moonPhase          precipIntensityMax
##          0          0
##          uvIndexTime          temperatureMin
##          0          0
##          temperatureMinTime          temperatureMax
##          0          0

```

```
##          temperatureMaxTime      apparentTemperatureMin
##                      0                      0
## apparentTemperatureMinTime      apparentTemperatureMax
##                      0                      0
## apparentTemperatureMaxTime
##                      0
```

```
cab_ds_distinct <- na.omit(cab_ds)
cab_ds_distinct<- cab_ds_distinct %>% distinct()
print(paste("The number of records removed : ", nrow(cab_ds) - nrow(cab_ds_distinct)))
```

```
## [1] "The number of records removed : 55095"
```

Exploratory Data Analysis

Who offers the most rides,Uber or lyft?

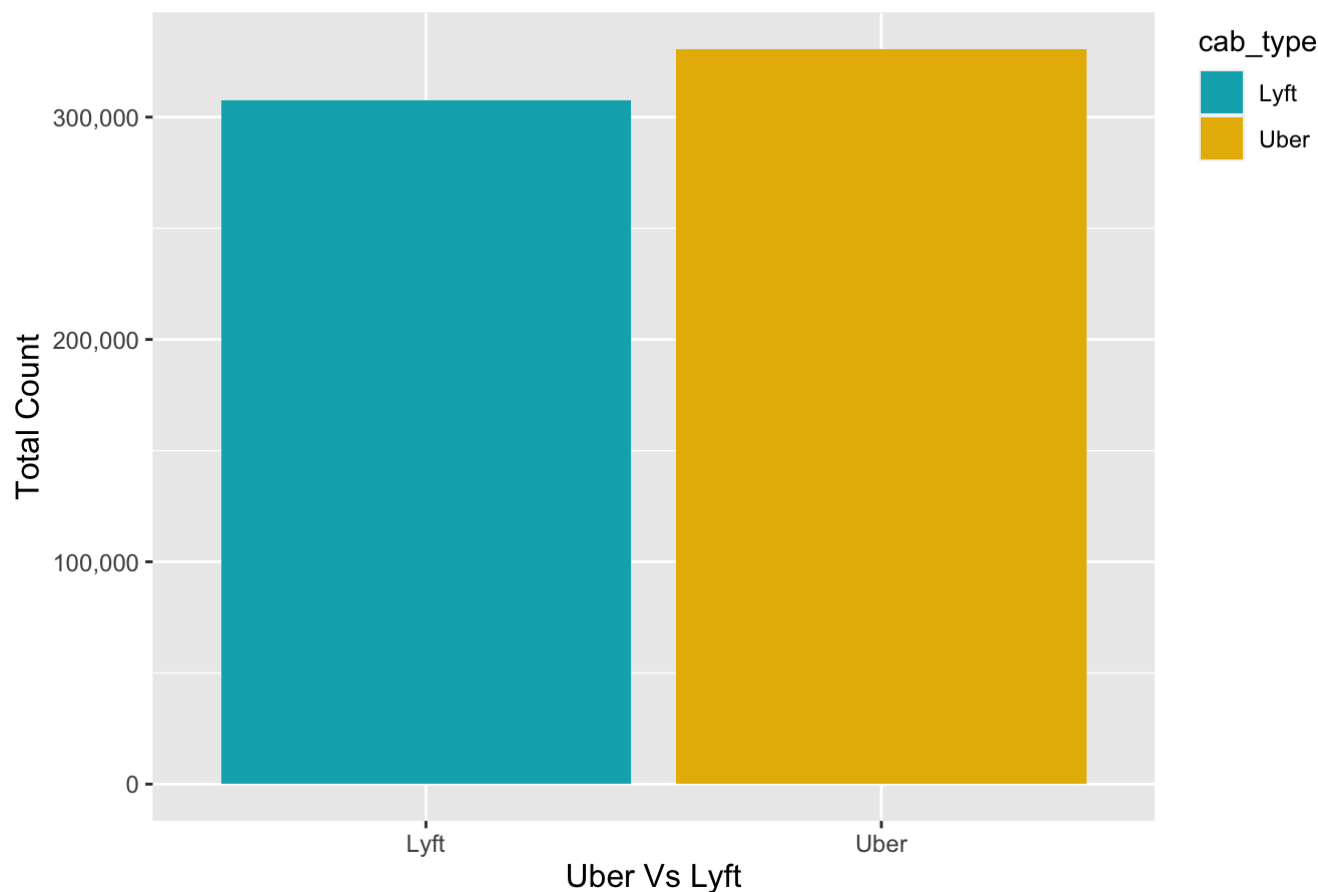
```
cab_ds_distinct %>% group_by(cab_type) %>%
  summarise("Total_Count" = length(id),
            'Percentage' = (length(id) / nrow(cab_ds_distinct)) * 100)
```

```
## # A tibble: 2 × 3
##   cab_type `Total_Count` Percentage
##   <chr>      <int>      <dbl>
## 1 Lyft      307408      48.2
## 2 Uber      330568      51.8
```

```
bp <-ggplot()+
  geom_bar(data=cab_ds_distinct,mapping=aes(x=cab_type, fill=cab_type))+
  scale_y_continuous(breaks = seq(0,1000000,100000),labels=scales::comma)+
  labs(x="Uber Vs Lyft",
       y="Total Count")+
  labs(title="Who offers the most rides")+
  theme(plot.title =element_text(hjust = 0.50,size=15),
        legend.justification = c("right", "top"),
        axis.title = element_text(size=12),
        axis.text = element_text(size=09))+
  theme(plot.caption=element_text(size=10))

bp + scale_fill_manual(values = c("#00AFBB", "#E7B800"))
```


Who offers the most rides



Lyft: Per Surge Multiplier - Total Rides vs Hour of the Day


```
surged_data <- cab_ds %>%
  filter(cab_type == "Lyft", surge_multiplier > 1.00) %>%
  dplyr::group_by(hour, surge_multiplier) %>%
  dplyr::summarize(total_rides = n())
```

```
## `summarise()` has grouped output by 'hour'. You can override using the
## `.groups` argument.
```

```
surged_data$surge_multiplier <- as.factor(surged_data$surge_multiplier)

lyft_surged_data <- ggplot(surged_data, aes(hour, total_rides, color = surge_multiplie
r)) +
  geom_point(alpha=0.8, size=1, aes(color = surge_multiplier)) +
  geom_line(aes(color = surge_multiplier)) + ggtitle("Lyft: Per Surge Multiplier -
Total Rides vs Hour of the Day") +
  facet_wrap(~surge_multiplier, ncol=1, scales="free") + xlab("Hour") + ylab("Tota
l Rides") +
  guides(color=guide_legend(ncol=1)) + theme(legend.position="none",
                                             panel.border = element_blank(),
                                             panel.spacing.x = unit(0,"line"))
lyft_surged_data
```

Lyft: Per Surge Multiplier - Total Rides vs Hour of the Day



Minimum and maximum fare prices

```
df<-sqldf("select source ,destination, cab_type ,avg(price) as average_price,min(price)
as minimum_price,max(price) as maximum_price from cab_ds group by source, destination,c
ab_type order by cab_type")
CrossTable(cab_ds$surge_multiplier, cab_ds$cab_type)
```

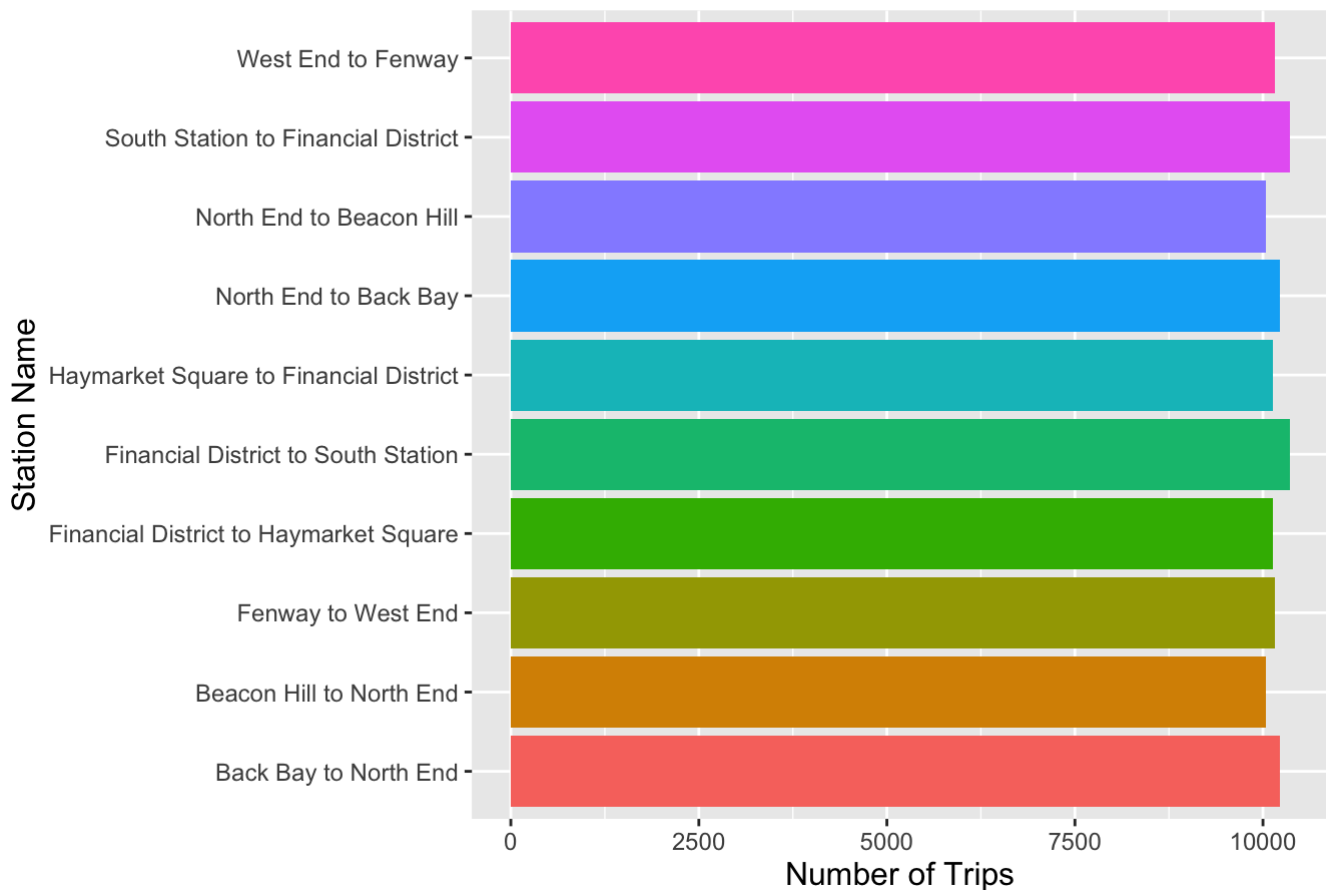
```
##
##
##      Cell Contents
## |-----|
## |                      N |
## | Chi-square contribution |
## |      N / Row Total |
## |      N / Col Total |
## |      N / Table Total |
## |-----|
##
##
## Total Observations in Table:  693071
##
##
##      cab_ds$cab_type
## cab_ds$surge_multiplier |      Lyft |      Uber | Row Total |
## -----|-----|-----|-----|
##              1 |      286433 |      385663 |      672096 |
##              |      456.978 |      364.253 |              |
##              |      0.426 |      0.574 |      0.970 |
##              |      0.932 |      1.000 |              |
##              |      0.413 |      0.556 |              |
## -----|-----|-----|-----|
##              1.25 |      11085 |           0 |      11085 |
##              |      7738.535 |      6168.307 |              |
##              |      1.000 |      0.000 |      0.016 |
##              |      0.036 |      0.000 |              |
##              |      0.016 |      0.000 |              |
## -----|-----|-----|-----|
##              1.5 |      5065 |           0 |      5065 |
##              |      3535.921 |      2818.446 |              |
##              |      1.000 |      0.000 |      0.007 |
##              |      0.016 |      0.000 |              |
##              |      0.007 |      0.000 |              |
## -----|-----|-----|-----|
##              1.75 |      2420 |           0 |      2420 |
##              |      1689.423 |      1346.622 |              |
##              |      1.000 |      0.000 |      0.003 |
##              |      0.008 |      0.000 |              |
##              |      0.003 |      0.000 |              |
## -----|-----|-----|-----|
##              2 |      2239 |           0 |      2239 |
##              |      1563.065 |      1245.903 |              |
##              |      1.000 |      0.000 |      0.003 |
##              |      0.007 |      0.000 |              |
##              |      0.003 |      0.000 |              |
## -----|-----|-----|-----|
##              2.5 |      154 |           0 |      154 |
##              |      107.509 |      85.694 |              |
##              |      1.000 |      0.000 |      0.000 |
##              |      0.001 |      0.000 |              |
```

```
##          |          0.000 |          0.000 |          |
## -----|-----|-----|-----|
##          3 |          12 |          0 |          12 |
##          |          8.377 |          6.677 |          |
##          |          1.000 |          0.000 |          0.000 |
##          |          0.000 |          0.000 |          |
##          |          0.000 |          0.000 |          |
## -----|-----|-----|-----|
##          Column Total |          307408 |          385663 |          693071 |
##          |          0.444 |          0.556 |          |
## -----|-----|-----|-----|
##
##
```

Top 10 most Popular Stations

```
popular_station<-sqldf("select source,destination,(source|| ' to ' ||destination) as sta
tion_name,count(id) as number_of_trips from cab_ds group by source,destination order by
count(id) desc LIMIT 10")
ggplot(data=popular_station,aes(x=number_of_trips, y=station_name, fill=station_name))+g
eom_bar(stat='identity')+
  labs(x="Number of Trips",y="Station Name")+
  labs(title=" Top 10 Popular Stations ")+
  theme(plot.title =element_text(hjust = 0.5,size=15),
        legend.position = c(2.50, .50),
        legend.justification = c("right", "top"),
        axis.title = element_text(size=12),
        axis.text = element_text(size=09))+
  theme(plot.caption=element_text(size=10))
```

Top 10 Popular Stations



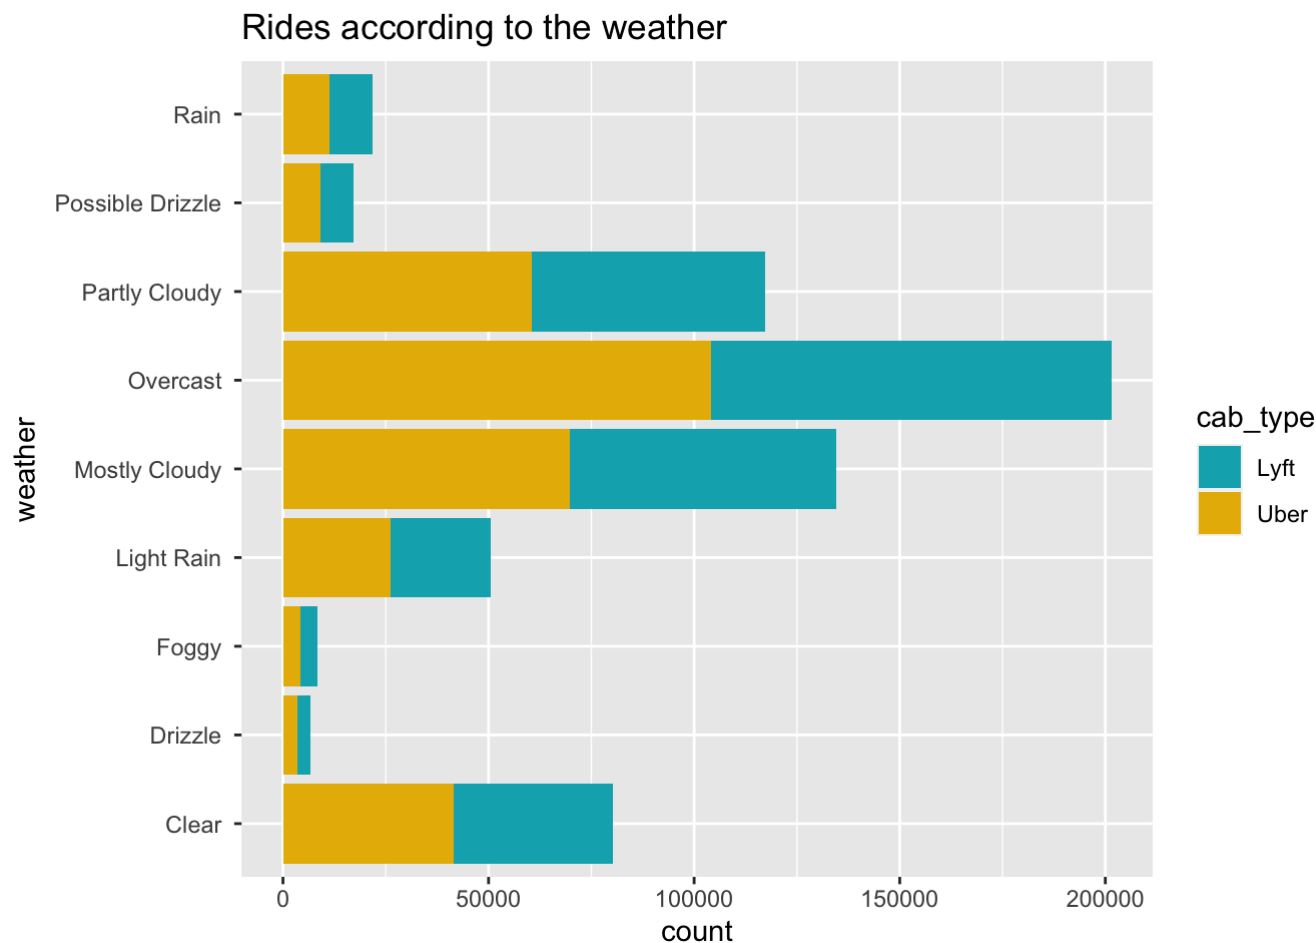
Weather affects the rides

```
cab_ds_distinct %>% group_by(short_summary) %>%
  summarise(count = length(id), 'Percentage' = (length(id) / nrow(cab_ds_distinct)) * 100)
```

```
## # A tibble: 9 × 3
##   short_summary      count Percentage
##   <chr>             <int>      <dbl>
## 1 " Clear "         80256      12.6
## 2 " Drizzle "       6725       1.05
## 3 " Foggy "        8292       1.30
## 4 " Light Rain "   50488       7.91
## 5 " Mostly Cloudy " 134603      21.1
## 6 " Overcast "     201429      31.6
## 7 " Partly Cloudy " 117226      18.4
## 8 " Possible Drizzle " 17176       2.69
## 9 " Rain "         21781       3.41
```

```
bp <- cab_ds_distinct %>%
  ggplot(aes(short_summary, fill=cab_type)) +
  labs(x="weather", title="Rides according to the weather") +
  geom_bar()+ coord_flip()

bp + scale_fill_manual(values = c("#00AFBB", "#E7B800"))
```




Temperature affects the ride's price

```
df2<-sqldf("select temperature, price , cab_type from cab_ds group by cab_type,temperature")
bp <- df2 %>%
  ggplot(aes(temperature, fill=cab_type)) +
  labs(x="Temperature", title="Cabs affected due to temperature") +
  geom_histogram()
bp + scale_fill_manual(values = c("#00AFBB", "#E7B800"))
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

Cabs affected due to temperature




Whether the passengers opt for cabs

```
document_tm <- Corpus(VectorSource(cab_ds$long_summary))
mat <- as.matrix(TermDocumentMatrix(document_tm))
vec <- sort(rowSums(mat), decreasing = TRUE)
word_corpus <- data.frame(word = names(vec), freq = vec)
set.seed(3)
wordcloud(word_corpus$word, freq = word_corpus$freq, colors = brewer.pal(8, "Dark2"))
```



file:///Users/Shared/pranjalnaik/Pranjal DPA/Src/DPA_Final_Project.html



Price range between Uber and Lyft

```
lyft<-sqldf("select * from cab_ds where cab_type='Lyft'")
uber<-sqldf("select * from cab_ds where cab_type='Uber'")
summary(lyft$price)
```


```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      2.50    9.00   16.50   17.35   22.50   97.50
```

```
summary(uber$price)
```


```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.   NA's
##      4.5    9.0    12.5    15.8   21.5    89.5  55095
```

```
hist(uber$price, col = "blue", density = 50, angle = 135, breaks = 40, xlim = c(0,80), m
ain = "Histogram of Uber & Lyft price")
hist(lyft$price, col = "green", density = 50, add = TRUE, breaks = 40)
```

Histogram of Uber & Lyft price



```
boxplot(cab_ds$price~cab_ds$cab_type,xlab='price', ylab='cab_type', data= cab_ds, horizontal = TRUE)
```




Heatmap for specific location and hours

```
bt<-cab_ds %>% select(price,cab_type,name,distance,short_summary,hour,source,destination) %>% filter(name!="WAV") %>% filter(name!="Lux") %>% filter(price>=0)
bt$name_f<-factor(bt$name,
                  levels=c("UberPool","Shared","UberX","Lyft","UberXL","Lyft XL",
                           "Black","Lux Black","Black SUV","Lux Black XL"))
levels(bt$name_f) <- list("Share" = c("UberPool","Shared"),
                        "Normal" = c("UberX","Lyft"),
                        "SUV" = c("UberXL","Lyft XL"),
                        "Lux" = c("Black","Lux Black"),
                        "Lux SUV" = c("Black SUV","Lux Black XL"))

bt<-bt %>% select(price,cab_type,name,name_f,distance,short_summary,hour,source,destination) %>% filter(name!="WAV") %>% filter(name!="Lux") %>% filter(price>=0)
bt1<-bt %>% select(price,cab_type,name_f,hour,source, destination) %>% filter(destination=="Northeastern University") %>% filter(source=="Theatre District") %>% filter(price>=0)
ggplot(bt1, aes(name_f,hour ))+
  geom_raster(aes(fill = price))+
  scale_fill_gradientn(colours=c("red","yellow"),name="Price")+
  labs(title = "Uber VS Lyft: Heat Map for Product types and Hours", x = "Product types",
       y = "Hours")+
  theme_bw()+facet_wrap(~cab_type)
```

Uber VS Lyft: Heat Map for Product types and Hours



Data Modelling

Loading pre processed Data and factoring required columns

Split data to train and test

```

wd <- weekdays(as.POSIXlt(cab_ds$datetime), abbreviate = TRUE)

cab_ds['Fri'] = as.integer(wd=='Fri')
cab_ds['Sat'] = as.integer(wd=='Sat')
cab_ds['Sun'] = as.integer(wd=='Sun')

#change short Summary of weather to binary variables
ss_data <- unique(cab_ds$short_summary)
for (i in ss_data)
{
  cab_ds[i] = as.integer(cab_ds$name == i)
}

for (p in unique(cab_ds$name))
{
  cab_ds[p] = as.integer(cab_ds$name == p)
}

lyft<-sqldf("select [distance],[surge_multiplier],[Fri], [Sat],[Sun],[Shared],[Lyft XL],
[Lux Black XL], [LUX],[Lux Black],[ Mostly Cloudy ], [ Rain ], [ Partly Cloudy ],[ Overc
ast ], [ Light Rain ], [ Foggy ], [ Possible Drizzle ],[ Drizzle ], price from cab_ds wh
ere cab_type='Lyft'")
uber<-sqldf("select [distance],[surge_multiplier],[Fri], [Sat],[Sun],[UberPool],[UberX
L],[Black],[Black SUV], [WAV],[ Mostly Cloudy ], [ Rain ], [ Partly Cloudy ],[ Overcast
 ], [ Light Rain ], [ Foggy ], [ Possible Drizzle ],[ Drizzle ], price from cab_ds where
cab_type='Uber'")

colnames(uber)[9] ="Black_SUV"
colnames(uber)[11] ="Mostly_Cloudy"
colnames(uber)[12] ="Rain"
colnames(uber)[13] ="Partly_Cloudy"
colnames(uber)[14] ="Overcast"
colnames(uber)[15] ="Light_Rain"
colnames(uber)[16] ="Foggy"
colnames(uber)[17] ="Possible_Drizzle"
colnames(uber)[18] ="Drizzle"

colnames(lyft)[7] ="Lyft_XL"
colnames(lyft)[8] ="Lux_Black_XL"
colnames(lyft)[10] ="Lux_Black"
colnames(lyft)[11] ="Mostly_Cloudy"
colnames(lyft)[12] ="Rain"
colnames(lyft)[13] ="Partly_Cloudy"
colnames(lyft)[14] ="Overcast"
colnames(lyft)[15] ="Light_Rain"
colnames(lyft)[16] ="Foggy"
colnames(lyft)[17] ="Possible_Drizzle"
colnames(lyft)[18] ="Drizzle"

#Uber

```

```
#selecting on numeric data
numeric_index = sapply(uber,is.numeric)
numeric_data = uber[,numeric_index]

#divide into train & test
train_index = sample(1:nrow(uber), 0.9 * nrow(uber))
uber_train = uber[train_index,]
uber_test = uber[-train_index,]

uber_train<-na.omit(uber_train)
sapply(uber_train, function(x) sum(is.na(x)))
```

```
##          distance surge_multiplier          Fri          Sat
##          0          0          0          0
##          Sun          UberPool          UberXL          Black
##          0          0          0          0
##          Black_SUV          WAV          Mostly_Cloudy          Rain
##          0          0          0          0
##          Partly_Cloudy          Overcast          Light_Rain          Foggy
##          0          0          0          0
## Possible_Drizzle          Drizzle          price
##          0          0          0
```

```
uber_test<-na.omit(uber_test)
sapply(uber_test, function(x) sum(is.na(x)))
```

```
##          distance surge_multiplier          Fri          Sat
##          0          0          0          0
##          Sun          UberPool          UberXL          Black
##          0          0          0          0
##          Black_SUV          WAV          Mostly_Cloudy          Rain
##          0          0          0          0
##          Partly_Cloudy          Overcast          Light_Rain          Foggy
##          0          0          0          0
## Possible_Drizzle          Drizzle          price
##          0          0          0
```

```
#lyft
#selecting on numeric data
numeric_index = sapply(lyft,is.numeric)
numeric_data = uber[,numeric_index]

#divide into train & test
train_index = sample(1:nrow(lyft), 0.9 * nrow(lyft))
lyft_train = lyft[train_index,]
lyft_test = lyft[-train_index,]

lyft_train<-na.omit(lyft_train)
sapply(lyft_train, function(x) sum(is.na(x)))
```

```
##           distance surge_multiplier           Fri           Sat
##           0           0           0           0
##           Sun           Shared           Lyft_XL           Lux_Black_XL
##           0           0           0           0
##           Lux           Lux_Black           Mostly_Cloudy           Rain
##           0           0           0           0
##           Partly_Cloudy           Overcast           Light_Rain           Foggy
##           0           0           0           0
## Possible_Drizzle           Drizzle           price
##           0           0           0
```

```
lyft_test<-na.omit(lyft_test)
sapply(lyft_test, function(x) sum(is.na(x)))
```

```
##           distance surge_multiplier           Fri           Sat
##           0           0           0           0
##           Sun           Shared           Lyft_XL           Lux_Black_XL
##           0           0           0           0
##           Lux           Lux_Black           Mostly_Cloudy           Rain
##           0           0           0           0
##           Partly_Cloudy           Overcast           Light_Rain           Foggy
##           0           0           0           0
## Possible_Drizzle           Drizzle           price
##           0           0           0
```

Linear Regression



```
#Uber
uber_lm_model = lm(price ~., data = uber_train)
summary(uber_lm_model)
```



```
##
## Call:
## lm(formula = price ~ ., data = uber_train)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-10.648	-1.420	-0.283	1.074	54.892

```
##
## Coefficients: (9 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   4.407800   0.014150  311.500 <2e-16 ***
## distance      2.445569   0.003772  648.374 <2e-16 ***
## surge_multiplier      NA          NA      NA      NA
## Fri           -0.022536   0.013601   -1.657  0.0975 .
## Sat            0.012649   0.013599    0.930  0.3523
## Sun           -0.031185   0.013512   -2.308  0.0210 *
## UberPool      -1.003611   0.015404  -65.151 <2e-16 ***
## UberXL         5.915176   0.015413  383.776 <2e-16 ***
## Black         10.773591   0.015407  699.253 <2e-16 ***
## Black_SUV     20.529696   0.015424 1331.030 <2e-16 ***
## WAV            0.002248   0.015417    0.146  0.8841
## Mostly_Cloudy      NA          NA      NA      NA
## Rain            NA          NA      NA      NA
## Partly_Cloudy     NA          NA      NA      NA
## Overcast         NA          NA      NA      NA
## Light_Rain       NA          NA      NA      NA
## Foggy            NA          NA      NA      NA
## Possible_Drizzle  NA          NA      NA      NA
## Drizzle          NA          NA      NA      NA
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.425 on 297500 degrees of freedom
## Multiple R-squared:  0.9198, Adjusted R-squared:  0.9198
## F-statistic: 3.79e+05 on 9 and 297500 DF, p-value: < 2.2e-16
```

```
plot (uber_lm_model)
```



```
#prediction
uber_pred = predict(uber_lm_model, uber_test[,1:18])
```

```
## Warning in predict.lm(uber_lm_model, uber_test[, 1:18]): prediction from a rank-
## deficient fit may be misleading
```

```
#Correlation Matrix
actuals_predicts <- data.frame(cbind(actuals=uber_test$price, predicted=uber_pred))
correlation_accuracy <- cor(actuals_predicts)
correlation_accuracy
```

```
##           actuals predicted
## actuals    1.0000000  0.9591117
## predicted  0.9591117  1.0000000
```

```
#Evaluation
mat_lr_uber<- regr.eval(uber_test[,19], uber_pred)#, stats = c('mape', 'rmse'))
print(mat_lr_uber)
```

```
##           mae           mse           rmse           mape
## 1.6697108 5.8347082 2.4155141 0.1191017
```

```
errors = abs(uber_pred - uber_test$price)
mape = 100 * (errors / uber_test$price)
uber_lr_accuracy = 100 - mean(mape)
sprintf("The Accuracy of Linear Regression for Uber :%f",uber_lr_accuracy)
```



```
## [1] "The Accuracy of Linear Regression for Uber :88.089834"
```



```
#-----
-----
#-----
-----

#lyft
lyft_lm_model = lm(price ~., data = lyft_train)
summary(lyft_lm_model)
```

```
##
## Call:
## lm(formula = price ~ ., data = lyft_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -21.744  -1.417  -0.165   1.422  40.354
##
## Coefficients: (8 not defined because of singularities)
##              Estimate Std. Error  t value Pr(>|t|)
## (Intercept)  -1.640e+01  3.988e-02 -411.217  <2e-16 ***
## distance      3.243e+00  4.428e-03  732.461  <2e-16 ***
## surge_multiplier 1.822e+01  3.568e-02  510.849  <2e-16 ***
## Fri          -1.054e-02  1.467e-02   -0.719    0.472
## Sat           1.304e-02  1.465e-02    0.890    0.374
## Sun           7.566e-04  1.461e-02    0.052    0.959
## Shared       -2.895e+00  1.670e-02 -173.332  <2e-16 ***
## Lyft_XL       5.696e+00  1.666e-02  341.886  <2e-16 ***
## Lux_Black_XL  2.273e+01  1.666e-02 1363.977  <2e-16 ***
## Lux           8.171e+00  1.665e-02  490.736  <2e-16 ***
## Lux_Black     1.347e+01  1.665e-02  808.881  <2e-16 ***
## Mostly_Cloudy      NA          NA      NA      NA
## Rain             NA          NA      NA      NA
## Partly_Cloudy     NA          NA      NA      NA
## Overcast          NA          NA      NA      NA
## Light_Rain        NA          NA      NA      NA
## Foggy            NA          NA      NA      NA
## Possible_Drizzle   NA          NA      NA      NA
## Drizzle           NA          NA      NA      NA
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.529 on 276656 degrees of freedom
## Multiple R-squared:  0.9363, Adjusted R-squared:  0.9363
## F-statistic: 4.068e+05 on 10 and 276656 DF, p-value: < 2.2e-16
```

```
plot(lyft_lm_model)
```




```
#prediction
lyft_pred = predict(lyft_lm_model, lyft_test[,1:18])
```

```
## Warning in predict.lm(lyft_lm_model, lyft_test[, 1:18]): prediction from a rank-
## deficient fit may be misleading
```

```
#Correlation Matrix
actuals_predicts <- data.frame(cbind(actuals=lyft_test$price, predicted=lyft_pred))
correlation_accuracy <- cor(actuals_predicts)
correlation_accuracy
```

```
##           actuals predicted
## actuals    1.0000000  0.9683375
## predicted  0.9683375  1.0000000
```

```
#Evaluation
mat_lr_lyft<- regr.eval(lyft_test[,19], lyft_pred)#, stats = c('mape', 'rmse'))
print(mat_lr_lyft)
```

```
##           mae           mse           rmse           mape
## 1.8068009  6.2327136  2.4965403  0.1493054
```

```
errors = abs(lyft_pred - lyft_test$price)
mape = 100 * (errors / lyft_test$price)
lyft_lr_accuracy = 100 - mean(mape)
sprintf("The Accuracy of Linear Regression for Lyft :%f",lyft_lr_accuracy)
```

```
## [1] "The Accuracy of Linear Regression for Lyft :85.069462"
```

Decision Tree

```
#Uber
uber_rpart_model = rpart(price ~., data = uber_train, method="anova")
summary(uber_rpart_model)
```

```

## Call:
## rpart(formula = price ~ ., data = uber_train, method = "anova")
##   n= 297510
##
##           CP nsplit  rel error      xerror      xstd
## 1 0.57182252      0 1.00000000 1.00001009 0.0028723813
## 2 0.16623160      1 0.42817748 0.42818237 0.0016622376
## 3 0.06681990      2 0.26194589 0.26195252 0.0012964846
## 4 0.03144997      3 0.19512599 0.19513268 0.0010827101
## 5 0.02763747      4 0.16367602 0.16368297 0.0010058718
## 6 0.01931217      5 0.13603855 0.13604616 0.0008734574
## 7 0.01245846      6 0.11672639 0.11673341 0.0007865145
## 8 0.01026684      7 0.10426792 0.10427501 0.0007714845
## 9 0.01000000      8 0.09400108 0.09525306 0.0007232491
##
## Variable importance
## Black_SUV      Black  distance      UberXL
##           63         18          11          7
##
## Node number 1: 297510 observations,      complexity param=0.5718225
##   mean=15.7961, MSE=73.33058
##   left son=2 (248034 obs) right son=3 (49476 obs)
##   Primary splits:
##     Black_SUV < 0.5   to the left,  improve=0.57182250, (0 missing)
##     UberPool  < 0.5   to the right, improve=0.13580220, (0 missing)
##     WAV       < 0.5   to the right, improve=0.09929631, (0 missing)
##     distance  < 2.295 to the left,  improve=0.07445334, (0 missing)
##     Black     < 0.5   to the left,  improve=0.06162716, (0 missing)
##
## Node number 2: 248034 observations,      complexity param=0.1662316
##   mean=12.90399, MSE=32.96959
##   left son=4 (198344 obs) right son=5 (49690 obs)
##   Primary splits:
##     Black     < 0.5   to the left,  improve=0.44348150, (0 missing)
##     distance  < 2.195 to the left,  improve=0.13931610, (0 missing)
##     UberPool  < 0.5   to the right, improve=0.13112140, (0 missing)
##     WAV       < 0.5   to the right, improve=0.07483935, (0 missing)
##     UberXL    < 0.5   to the left,  improve=0.05861073, (0 missing)
##
## Node number 3: 49476 observations,      complexity param=0.02763747
##   mean=30.29488, MSE=23.52232
##   left son=6 (37701 obs) right son=7 (11775 obs)
##   Primary splits:
##     distance  < 2.865 to the left,  improve=5.18096e-01, (0 missing)
##     Sat       < 0.5   to the left,  improve=5.43687e-05, (0 missing)
##     Fri       < 0.5   to the right, improve=3.19882e-05, (0 missing)
##     Sun       < 0.5   to the left,  improve=6.88000e-06, (0 missing)
##
## Node number 4: 198344 observations,      complexity param=0.0668199
##   mean=10.99009, MSE=16.76863
##   left son=8 (148728 obs) right son=9 (49616 obs)
##   Primary splits:

```

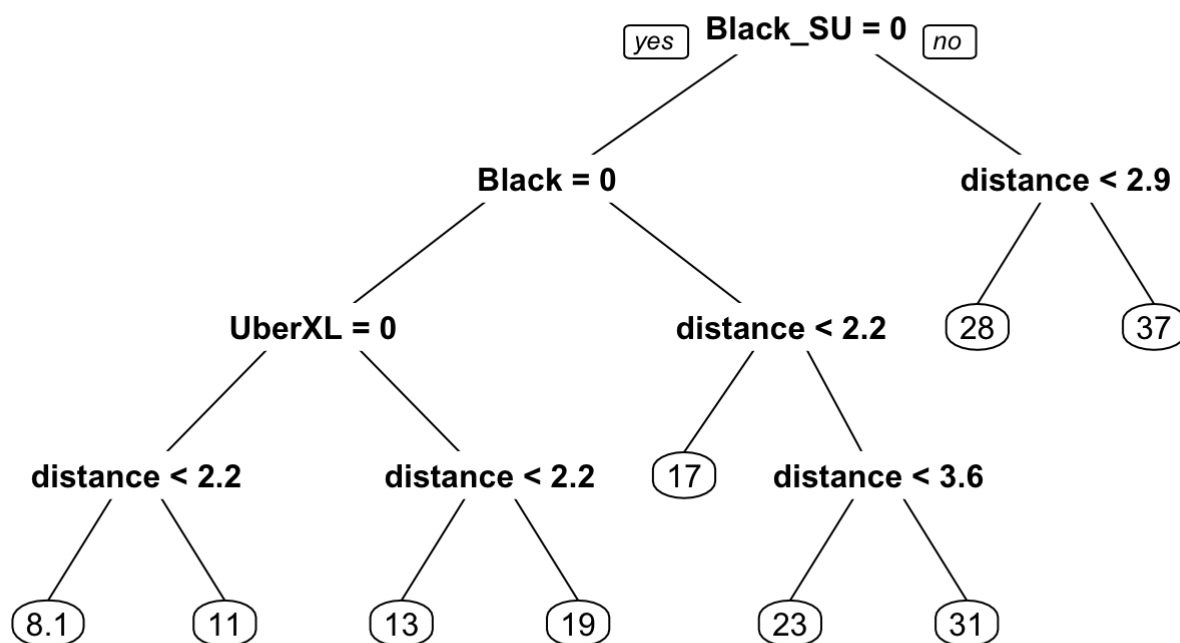
```
##      UberXL    < 0.5   to the left,  improve=4.383045e-01, (0 missing)
##      distance < 2.195 to the left,  improve=1.814092e-01, (0 missing)
##      UberPool < 0.5   to the right, improve=9.995433e-02, (0 missing)
##      WAV       < 0.5   to the right, improve=3.002805e-02, (0 missing)
##      Fri       < 0.5   to the right, improve=1.172281e-05, (0 missing)
##
## Node number 5: 49690 observations,      complexity param=0.03144997
##   mean=20.54357, MSE=24.65318
##   left son=10 (25181 obs) right son=11 (24509 obs)
##   Primary splits:
##     distance < 2.195 to the left,  improve=5.600991e-01, (0 missing)
##     Fri      < 0.5   to the right, improve=3.978983e-05, (0 missing)
##     Sat      < 0.5   to the right, improve=1.947026e-05, (0 missing)
##     Sun      < 0.5   to the left,  improve=1.001368e-06, (0 missing)
##
## Node number 6: 37701 observations
##   mean=28.34392, MSE=7.632643
##
## Node number 7: 11775 observations
##   mean=36.54144, MSE=23.19123
##
## Node number 8: 148728 observations,      complexity param=0.01245846
##   mean=9.424234, MSE=5.735292
##   left son=16 (75588 obs) right son=17 (73140 obs)
##   Primary splits:
##     distance < 2.195 to the left,  improve=3.186419e-01, (0 missing)
##     UberPool < 0.5   to the right, improve=3.957829e-02, (0 missing)
##     WAV      < 0.5   to the left,  improve=9.862602e-03, (0 missing)
##     Sat      < 0.5   to the right, improve=4.397628e-05, (0 missing)
##     Sun      < 0.5   to the left,  improve=2.553613e-05, (0 missing)
##   Surrogate splits:
##     Sun < 0.5   to the left,  agree=0.509, adj=0.001, (0 split)
##
## Node number 9: 49616 observations,      complexity param=0.01931217
##   mean=15.68386, MSE=20.46067
##   left son=18 (25185 obs) right son=19 (24431 obs)
##   Primary splits:
##     distance < 2.195 to the left,  improve=4.150268e-01, (0 missing)
##     Fri      < 0.5   to the right, improve=1.219591e-05, (0 missing)
##     Sat      < 0.5   to the left,  improve=7.710355e-06, (0 missing)
##     Sun      < 0.5   to the right, improve=1.881805e-06, (0 missing)
##
## Node number 10: 25181 observations
##   mean=16.87755, MSE=3.94707
##
## Node number 11: 24509 observations,      complexity param=0.01026684
##   mean=24.31011, MSE=17.93197
##   left son=22 (20484 obs) right son=23 (4025 obs)
##   Primary splits:
##     distance < 3.575 to the left,  improve=5.096473e-01, (0 missing)
##     Fri      < 0.5   to the right, improve=2.516569e-05, (0 missing)
##     Sat      < 0.5   to the right, improve=7.776256e-08, (0 missing)
```

```
##          Sun          < 0.5   to the left,  improve=7.515796e-09, (0 missing)
##
## Node number 16: 75588 observations
##   mean=8.094453, MSE=2.347904
##
## Node number 17: 73140 observations
##   mean=10.79852, MSE=5.51988
##
## Node number 18: 25185 observations
##   mean=12.81376, MSE=6.920754
##
## Node number 19: 24431 observations
##   mean=18.64254, MSE=17.17293
##
## Node number 22: 20484 observations
##   mean=22.97005, MSE=6.733298
##
## Node number 23: 4025 observations
##   mean=31.12994, MSE=19.27517
```

```
#identify best cp value to use
best <- uber_rpart_model$scptable[which.min(uber_rpart_model$scptable[, "xerror"]), "CP"]

#produce a pruned tree based on the best cp value
pruned_tree <- prune(uber_rpart_model, cp=best)

#plot the pruned tree
prp(pruned_tree)
```



```

#prediction
uber_pred_rpart = predict(uber_rpart_model, uber_test[,-19])

#Correlation Matrix
actuals_predicts <- data.frame(cbind(actuals=uber_test$price, predicted=uber_pred_rpart))
correlation_accuracy <- cor(actuals_predicts)
correlation_accuracy

```

```

##           actuals predicteds
## actuals    1.0000000  0.9520703
## predicted 0.9520703  1.0000000

```

```

#Evaluation
mat_dt_uber<- regr.eval(uber_test[,19], uber_pred_rpart)#, stats = c('mape','rmse'))
print(mat_dt_uber)

```

```

##           mae      mse      rmse      mape
## 1.7995829 6.8132767 2.6102254 0.1230596

```

```
errors = abs(uber_pred_rpart - uber_test$price)
mape = 100 * (errors / uber_test$price)
uber_dt_accuracy = 100 - mean(mape)
sprintf("The Accuracy of Decision Tree for Uber :%f",uber_dt_accuracy)
```

```
## [1] "The Accuracy of Decision Tree for Uber :87.694036"
```

```
#-----
-----
#-----
-----

# lyft
lyft_rpart_model = rpart(price ~., data = lyft_train, method="anova")
summary(lyft_rpart_model)
```

```

## Call:
## rpart(formula = price ~ ., data = lyft_train, method = "anova")
##   n= 276667
##
##           CP nsplit rel error   xerror   xstd
## 1  0.44585340     0 1.0000000 1.0000058 0.0035430159
## 2  0.15776524     1 0.5541466 0.5541521 0.0023072276
## 3  0.08367959     2 0.3963814 0.3963891 0.0019125673
## 4  0.03891566     3 0.3127018 0.3127098 0.0018211360
## 5  0.03231274     4 0.2737861 0.2738188 0.0017487835
## 6  0.03101511     5 0.2414734 0.2416787 0.0016521927
## 7  0.02459246     6 0.2104583 0.2106961 0.0014428580
## 8  0.02088014     7 0.1858658 0.1860868 0.0014020353
## 9  0.01746894     8 0.1649857 0.1651956 0.0013715584
## 10 0.01156837     9 0.1475167 0.1477901 0.0010403555
## 11 0.01103924    10 0.1359484 0.1353521 0.0008943156
## 12 0.01000000    11 0.1249091 0.1261764 0.0007299040
##
## Variable importance
##      Lux_Black_XL      Lux_Black      distance      Shared
##              51              18              12              10
## surge_multiplier              Lux      Lyft_XL
##              5              3              2
##
## Node number 1: 276667 observations,      complexity param=0.4458534
##   mean=17.34354, MSE=100.4256
##   left son=2 (230640 obs) right son=3 (46027 obs)
##   Primary splits:
##      Lux_Black_XL      < 0.5   to the left, improve=0.44585340, (0 missing)
##      Shared          < 0.5   to the right, improve=0.25535070, (0 missing)
##      distance        < 2.445 to the left, improve=0.09106581, (0 missing)
##      surge_multiplier < 1.125 to the left, improve=0.08077146, (0 missing)
##      Lux_Black       < 0.5   to the left, improve=0.06536165, (0 missing)
##
## Node number 2: 230640 observations,      complexity param=0.1577652
##   mean=14.35433, MSE=56.38975
##   left son=4 (184461 obs) right son=5 (46179 obs)
##   Primary splits:
##      Lux_Black      < 0.5   to the left, improve=0.33703760, (0 missing)
##      Shared         < 0.5   to the right, improve=0.30765310, (0 missing)
##      distance       < 2.175 to the left, improve=0.13826520, (0 missing)
##      surge_multiplier < 1.125 to the left, improve=0.09972451, (0 missing)
##      Lux            < 0.5   to the left, improve=0.05171951, (0 missing)
##
## Node number 3: 46027 observations,      complexity param=0.03101511
##   mean=32.32242, MSE=51.94622
##   left son=6 (28747 obs) right son=7 (17280 obs)
##   Primary splits:
##      distance      < 2.495 to the left, improve=3.604198e-01, (0 missing)
##      surge_multiplier < 1.375 to the left, improve=3.554192e-01, (0 missing)
##      Fri           < 0.5   to the right, improve=4.740028e-05, (0 missing)
##      Sat           < 0.5   to the left, improve=9.382979e-06, (0 missing)

```

```

##          Sun          < 0.5   to the left,  improve=4.275176e-07, (0 missing)
##
## Node number 4: 184461 observations,      complexity param=0.08367959
##   mean=12.17306, MSE=36.22639
##   left son=8 (46168 obs) right son=9 (138293 obs)
##   Primary splits:
##     Shared          < 0.5   to the right, improve=0.34792990, (0 missing)
##     Lux             < 0.5   to the left,  improve=0.28836720, (0 missing)
##     distance        < 1.955 to the left,  improve=0.15749800, (0 missing)
##     surge_multiplier < 1.125 to the left,  improve=0.11212680, (0 missing)
##     Lyft_XL         < 0.5   to the left,  improve=0.08942717, (0 missing)
##
## Node number 5: 46179 observations,      complexity param=0.03231274
##   mean=23.06736, MSE=42.00945
##   left son=10 (23531 obs) right son=11 (22648 obs)
##   Primary splits:
##     distance        < 2.175 to the left,  improve=4.627902e-01, (0 missing)
##     surge_multiplier < 1.375 to the left,  improve=2.271646e-01, (0 missing)
##     Sun             < 0.5   to the left,  improve=3.186539e-05, (0 missing)
##     Fri             < 0.5   to the left,  improve=9.812630e-06, (0 missing)
##     Sat             < 0.5   to the right, improve=1.424171e-06, (0 missing)
##   Surrogate splits:
##     surge_multiplier < 1.125 to the left,  agree=0.521, adj=0.023, (0 split)
##     Sun             < 0.5   to the left,  agree=0.510, adj=0.000, (0 split)
##
## Node number 6: 28747 observations,      complexity param=0.01156837
##   mean=28.9677, MSE=19.8281
##   left son=12 (27781 obs) right son=13 (966 obs)
##   Primary splits:
##     surge_multiplier < 1.375 to the left,  improve=5.638976e-01, (0 missing)
##     distance        < 1.955 to the left,  improve=1.489673e-01, (0 missing)
##     Fri             < 0.5   to the right, improve=3.106035e-04, (0 missing)
##     Sun             < 0.5   to the right, improve=2.727286e-06, (0 missing)
##     Sat             < 0.5   to the right, improve=7.353013e-07, (0 missing)
##
## Node number 7: 17280 observations,      complexity param=0.01746894
##   mean=37.90334, MSE=55.50882
##   left son=14 (16450 obs) right son=15 (830 obs)
##   Primary splits:
##     surge_multiplier < 1.375 to the left,  improve=5.060141e-01, (0 missing)
##     distance        < 3.555 to the left,  improve=2.237248e-01, (0 missing)
##     Sun             < 0.5   to the right, improve=5.680013e-05, (0 missing)
##     Sat             < 0.5   to the right, improve=7.306841e-06, (0 missing)
##     Fri             < 0.5   to the left,  improve=6.602096e-06, (0 missing)
##
## Node number 8: 46168 observations
##   mean=6.028537, MSE=4.441998
##
## Node number 9: 138293 observations,      complexity param=0.03891566
##   mean=14.22436, MSE=30.02528
##   left son=18 (68334 obs) right son=19 (69959 obs)
##   Primary splits:

```



```

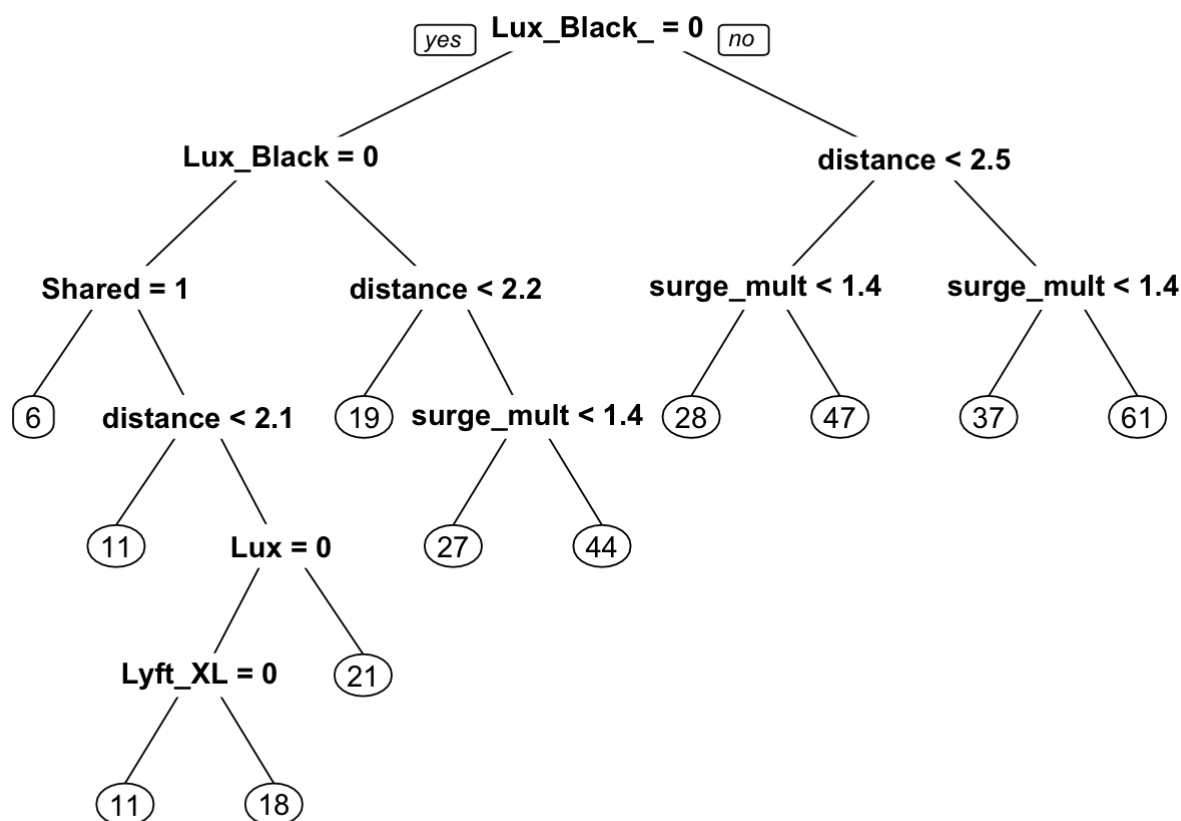
##      distance      < 2.125 to the left,  improve=2.603988e-01, (0 missing)
##      Lux           < 0.5   to the left,  improve=2.095525e-01, (0 missing)
##      surge_multiplier < 1.375 to the left,  improve=1.093318e-01, (0 missing)
##      Lyft_XL       < 0.5   to the left,  improve=1.901981e-02, (0 missing)
##      Sat           < 0.5   to the right, improve=9.215372e-06, (0 missing)
##  Surrogate splits:
##      surge_multiplier < 1.125 to the left,  agree=0.507, adj=0.002, (0 split)
##
## Node number 10: 23531 observations
##   mean=18.74162, MSE=10.94892
##
## Node number 11: 22648 observations,      complexity param=0.01103924
##   mean=27.56175, MSE=34.63985
##   left son=22 (21586 obs) right son=23 (1062 obs)
##   Primary splits:
##       surge_multiplier < 1.375 to the left,  improve=3.909626e-01, (0 missing)
##       distance         < 3.545 to the left,  improve=3.072661e-01, (0 missing)
##       Fri              < 0.5   to the left,  improve=5.558182e-05, (0 missing)
##       Sat              < 0.5   to the right, improve=1.157420e-05, (0 missing)
##       Sun              < 0.5   to the left,  improve=9.455698e-08, (0 missing)
##
## Node number 12: 27781 observations
##   mean=28.34417, MSE=7.205983
##
## Node number 13: 966 observations
##   mean=46.89959, MSE=50.09137
##
## Node number 14: 16450 observations
##   mean=36.71287, MSE=23.00623
##
## Node number 15: 830 observations
##   mean=61.49759, MSE=114.9096
##
## Node number 18: 68334 observations
##   mean=11.39514, MSE=11.97631
##
## Node number 19: 69959 observations,      complexity param=0.02459246
##   mean=16.98786, MSE=32.19954
##   left son=38 (46591 obs) right son=39 (23368 obs)
##   Primary splits:
##       Lux              < 0.5   to the left,  improve=0.3033266000, (0 missing)
##       surge_multiplier < 1.375 to the left,  improve=0.1417080000, (0 missing)
##       distance         < 3.545 to the left,  improve=0.1207882000, (0 missing)
##       Lyft_XL          < 0.5   to the left,  improve=0.0269986000, (0 missing)
##       Sat              < 0.5   to the right, improve=0.0000834332, (0 missing)
##   Surrogate splits:
##       Lyft_XL < 0.5   to the right, agree=0.666, adj=0.001, (0 split)
##
## Node number 22: 21586 observations
##   mean=26.74548, MSE=17.96197
##
## Node number 23: 1062 observations

```

```
## mean=44.15301, MSE=84.81816
##
## Node number 38: 46591 observations, complexity param=0.02088014
## mean=14.77456, MSE=23.21915
## left son=76 (23334 obs) right son=77 (23257 obs)
## Primary splits:
## Lyft_XL < 0.5 to the left, improve=5.362742e-01, (0 missing)
## surge_multiplier < 1.375 to the left, improve=1.488868e-01, (0 missing)
## distance < 3.545 to the left, improve=1.257325e-01, (0 missing)
## Sat < 0.5 to the right, improve=8.349234e-05, (0 missing)
## Fri < 0.5 to the right, improve=5.437046e-06, (0 missing)
## Surrogate splits:
## distance < 3.205 to the left, agree=0.502, adj=0.002, (0 split)
## Sat < 0.5 to the right, agree=0.502, adj=0.002, (0 split)
##
## Node number 39: 23368 observations
## mean=21.40072, MSE=20.86427
##
## Node number 76: 23334 observations
## mean=11.25167, MSE=5.12571
##
## Node number 77: 23257 observations
## mean=18.30911, MSE=16.4276
```

```
#identify best cp value to use
best <- lyft_rpart_model$cpstable[which.min(lyft_rpart_model$cpstable[, "xerror"]), "CP"]
#produce a pruned tree based on the best cp value
pruned_tree <- prune(lyft_rpart_model, cp=best)

#plot the pruned tree
prp(pruned_tree)
```



```
#prediction
```

```
lyft_pred_rpart = predict(lyft_rpart_model, lyft_test[,-19])
```

```
#Correlation Matrix
```

```
actuals_predicts <- data.frame(cbind(actuals=lyft_test$price, predicted=lyft_pred_rpart))
```

```
correlation_accuracy <- cor(actuals_predicts)
```

```
correlation_accuracy
```

```
##          actuals predicted
```

```
## actuals    1.0000000  0.9348973
```

```
## predicted 0.9348973  1.0000000
```

```
#Evaluation
```

```
mat_dt_lyft<- regr.eval(lyft_test[,19], lyft_pred_rpart)#, stats = c('mape','rmse'))
```

```
print(mat_dt_lyft)
```

```
##          mae          mse          rmse          mape
```

```
##  2.6463307 12.5978293  3.5493421  0.1969981
```

```
errors = abs(lyft_pred_rpart - lyft_test$price)
mape = 100 * (errors / lyft_test$price)
lyft_dt_accuracy = 100 - mean(mape)
sprintf("The Accuracy of Decision Tree for Lyft :%f",lyft_dt_accuracy)
```

```
## [1] "The Accuracy of Decision Tree for Lyft :80.300189"
```

Random Forest

```
#Uber
#head(uber_train)
uber_rmforest_model = randomForest(price ~., data = uber_train, importance = TRUE, ntree
= 100)
summary(uber_rmforest_model)
```

```
##           Length Class  Mode
## call              5 -none- call
## type              1 -none- character
## predicted        297510 -none- numeric
## mse               100 -none- numeric
## rsq               100 -none- numeric
## oob.times         297510 -none- numeric
## importance         36 -none- numeric
## importanceSD       18 -none- numeric
## localImportance    0 -none- NULL
## proximity          0 -none- NULL
## ntree              1 -none- numeric
## mtry               1 -none- numeric
## forest            11 -none- list
## coefs              0 -none- NULL
## y                 297510 -none- numeric
## test              0 -none- NULL
## inbag              0 -none- NULL
## terms              3 terms  call
```

```
#prediction
uber_pred_rmforest = predict(uber_rmforest_model, uber_test[,-19])

#Correlation Matrix
actuals_predicts <- data.frame(cbind(actuals=uber_test$price, predicted=uber_pred_rmfor
est))
correlation_accuracy <- cor(actuals_predicts)
correlation_accuracy
```

```
##           actuals predicteds
## actuals      1.0000000  0.9576262
## predicteds  0.9576262  1.0000000
```

```
#Evaluation
mat_rf_uber<- regr.eval(uber_test[,19], uber_pred_rmforest)#, stats = c('mape','rmse'))
print(mat_rf_uber)
```

```
##          mae          mse          rmse          mape
## 1.7504893 6.6244389 2.5737985 0.1288409
```

```
errors = abs(uber_pred_rmforest - uber_test$price)
mape = 100 * (errors / uber_test$price)
uber_rf_accuracy = 100 - mean(mape)
sprintf("The Accuracy of Random Forest for Uber :%f",uber_rf_accuracy)
```

```
## [1] "The Accuracy of Random Forest for Uber :87.115908"
```

```
#-----
#-----
#-----
#-----

lyft_rmforest_model = randomForest(price ~., data = lyft_train, importance = TRUE, ntree
= 100)
summary(lyft_rmforest_model)
```

```
##          Length Class  Mode
## call          5 -none- call
## type          1 -none- character
## predicted     276667 -none- numeric
## mse           100 -none- numeric
## rsq           100 -none- numeric
## oob.times     276667 -none- numeric
## importance     36 -none- numeric
## importanceSD   18 -none- numeric
## localImportance 0 -none- NULL
## proximity      0 -none- NULL
## ntree          1 -none- numeric
## mtry           1 -none- numeric
## forest        11 -none- list
## coefs          0 -none- NULL
## y             276667 -none- numeric
## test          0 -none- NULL
## inbag          0 -none- NULL
## terms         3 terms  call
```

```
#prediction
lyft_pred_rmforest = predict(lyft_rmforest_model, lyft_test[,-19])

#Correlation Matrix
actuals_predicts <- data.frame(cbind(actuals=lyft_test$price, predicted=lyft_pred_rmforest))
correlation_accuracy <- cor(actuals_predicts)
correlation_accuracy
```

```
##           actuals predicted
## actuals    1.0000000  0.9766981
## predicted 0.9766981  1.0000000
```

```
#Evaluation
mat_rf_lyft<- regr.eval(lyft_test[,19], lyft_pred_rmforest)#, stats = c('mape','rmse'))
print(mat_rf_lyft)
```

```
##           mae           mse           rmse           mape
## 1.7126200 5.4800038 2.3409408 0.1326613
```

```
errors = abs(lyft_pred_rmforest - lyft_test$price)
mape = 100 * (errors / lyft_test$price)
lyft_rf_accuracy = 100 - mean(mape)
sprintf("The Accuracy of Random Forest for Lyft :%f",lyft_rf_accuracy)
```

```
## [1] "The Accuracy of Random Forest for Lyft :86.733867"
```

Model Evaluation

```
# Uber
print("*****Uber Statistics*****")
```

```
## [1] "*****Uber Statistics*****"
```

```
tab <- matrix(c(mat_lr_uber["mae"],mat_dt_uber["mae"],mat_rf_uber["mae"],mat_lr_uber["mse"],mat_dt_uber["mse"],mat_rf_uber["mse"],mat_lr_uber["rmse"],mat_dt_uber["rmse"],mat_rf_uber["rmse"],mat_lr_uber["mape"],mat_dt_uber["mape"],mat_rf_uber["mape"], uber_lr_accuracy,uber_dt_accuracy,uber_rf_accuracy), ncol=3, byrow=TRUE)
colnames(tab) <- c("Linear Regression",'Decision Tree','Random Forest')
rownames(tab) <- c('MAE','MSE','RMSE','MAPE','Accuracy')
uber_tab <- as.table(tab)
uber_tab
```

```
##           Linear Regression Decision Tree Random Forest
## MAE           1.6697108       1.7995829       1.7504893
## MSE           5.8347082       6.8132767       6.6244389
## RMSE          2.4155141       2.6102254       2.5737985
## MAPE          0.1191017       0.1230596       0.1288409
## Accuracy      88.0898340      87.6940363      87.1159081
```

```
print("*****lyft Statitics*****")
```

```
## [1] "*****lyft Statitics*****"
```

```
# Lyft
tab <- matrix(c(mat_lr_lyft["mae"],mat_dt_lyft["mae"],mat_rf_lyft["mae"],mat_lr_lyft["mse"],mat_dt_lyft["mse"],mat_rf_lyft["mse"],mat_lr_lyft["rmse"],mat_dt_lyft["rmse"],mat_rf_lyft["rmse"],mat_lr_lyft["mape"],mat_dt_lyft["mape"],mat_rf_lyft["mape"], lyft_lr_accuracy,lyft_dt_accuracy,lyft_rf_accuracy), ncol=3, byrow=TRUE)
colnames(tab) <- c("Linear Regression",'Decision Tree','Random Forest')
rownames(tab) <- c('MAE','MSE','RMSE','MAPE',"Accuracy")
lyft_tab <- as.table(tab)
lyft_tab
```

```
##           Linear Regression Decision Tree Random Forest
## MAE           1.8068009       2.6463307       1.7126200
## MSE           6.2327136      12.5978293       5.4800038
## RMSE          2.4965403       3.5493421       2.3409408
## MAPE          0.1493054       0.1969981       0.1326613
## Accuracy      85.0694620      80.3001895      86.7338674
```

Based on different metrics, we can conclude that the Random Forest model is the best model and thus we could consider this as our final model.